

```
In [42]: import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_e
from sklearn.linear_model import LinearRegression
```

```
In [43]: import pandas as pd
from sklearn.preprocessing import StandardScaler
import numpy as np
```

```
In [44]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_e
```

```
In [45]: #A) DATA PROCESSING:

# 1) Load train & test data
train = pd.read_csv("california_housing_train.csv")
test = pd.read_csv("california_housing_test.csv")

print("Train shape:", train.shape)
print("Test shape:", test.shape)
X_train = train.drop(columns=['median_house_value'])
y_train = train['median_house_value']

X_test = test.drop(columns=['median_house_value'])
y_test = test['median_house_value']

# 2) Standardize features (zero mean, unit variance)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
X_train_scaled = pd.DataFrame(X_train_scaled, columns=X_train.columns)
X_test_scaled = pd.DataFrame(X_test_scaled, columns=X_test.columns)

# 3) Add intercept (bias) column of ones
X_train_scaled.insert(0, 'intercept', 1)
X_test_scaled.insert(0, 'intercept', 1)

print("Processed Train Features Shape:", X_train_scaled.shape)
print("Processed Test Features Shape:", X_test_scaled.shape)
```

Train shape: (17000, 9)

Test shape: (3000, 9)

Processed Train Features Shape: (17000, 9)

Processed Test Features Shape: (3000, 9)

```
In [46]: train.head()
```

Out [46]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	populati
0	-114.31	34.19	15.0	5612.0	1283.0	101
1	-114.47	34.40	19.0	7650.0	1901.0	112
2	-114.56	33.69	17.0	720.0	174.0	33
3	-114.57	33.64	14.0	1501.0	337.0	51
4	-114.57	33.57	20.0	1454.0	326.0	62

In [47]:

```
# === Dataset Information ===  
print("Info:")  
print(train.info())  
print("\nMissing Values:\n", train.isnull().sum())  
print("\nSummary Stats:\n", train.describe())
```

Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 17000 entries, 0 to 16999

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	longitude	17000 non-null	float64
1	latitude	17000 non-null	float64
2	housing_median_age	17000 non-null	float64
3	total_rooms	17000 non-null	float64
4	total_bedrooms	17000 non-null	float64
5	population	17000 non-null	float64
6	households	17000 non-null	float64
7	median_income	17000 non-null	float64
8	median_house_value	17000 non-null	float64

dtypes: float64(9)

memory usage: 1.2 MB

None

Missing Values:

longitude	0
latitude	0
housing_median_age	0
total_rooms	0
total_bedrooms	0
population	0
households	0
median_income	0
median_house_value	0

dtype: int64

Summary Stats:

	longitude	latitude	housing_median_age	total_rooms \
count	17000.000000	17000.000000	17000.000000	17000.000000
mean	-119.562108	35.625225	28.589353	2643.664412
std	2.005166	2.137340	12.586937	2179.947071
min	-124.350000	32.540000	1.000000	2.000000
25%	-121.790000	33.930000	18.000000	1462.000000
50%	-118.490000	34.250000	29.000000	2127.000000
75%	-118.000000	37.720000	37.000000	3151.250000
max	-114.310000	41.950000	52.000000	37937.000000


	total_bedrooms	population	households	median_income \
count	17000.000000	17000.000000	17000.000000	17000.000000
mean	539.410824	1429.573941	501.221941	3.883578
std	421.499452	1147.852959	384.520841	1.908157
min	1.000000	3.000000	1.000000	0.499900
25%	297.000000	790.000000	282.000000	2.566375
50%	434.000000	1167.000000	409.000000	3.544600
75%	648.250000	1721.000000	605.250000	4.767000
max	6445.000000	35682.000000	6082.000000	15.000100

	median_house_value
count	17000.000000
mean	207300.912353
std	115983.764387
min	14999.000000
25%	119400.000000
50%	180400.000000


```
75%      265000.000000
max      500001.000000
```

```
In [48]: # === Target Distribution ===
TARGET_NAME = train.columns[-1] # assume last col is target

plt.figure(figsize=(7,5))
sns.histplot(train[TARGET_NAME], kde=True, bins=40)
plt.title("Target Distribution (Before Transformation)")
plt.show()
```



```
In [49]: # === Resolve Right Skewness ===
fig, axes = plt.subplots(1,3, figsize=(15,4))
sns.histplot(np.log1p(train[TARGET_NAME]), kde=True, ax=axes[0]); axes[0].
sns.histplot(np.sqrt(train[TARGET_NAME]), kde=True, ax=axes[1]); axes[1].
sns.histplot(1/(train[TARGET_NAME]+1e-6), kde=True, ax=axes[2]); axes[2].
plt.show()
```




```
In [50]: # === Predictor Distribution + Outlier Removal ===
col = train.columns[0] # just pick first feature for demo, adjust as nee


plt.figure(figsize=(7,5))
```

```
sns.histplot(train[col], kde=True)
plt.title(f"Distribution of {col} (Before Outlier Removal)")
plt.show()


# IQR method
Q1, Q3 = train[col].quantile([0.25,0.75])
IQR = Q3 - Q1
low, high = Q1 - 1.5*IQR, Q3 + 1.5*IQR
train = train[(train[col] >= low) & (train[col] <= high)]

plt.figure(figsize=(7,5))
sns.histplot(train[col], kde=True)
plt.title(f"Distribution of {col} (After Outlier Removal)")
plt.show()
```





```
In [51]: # === Correlation Heatmap ===  
plt.figure(figsize=(10,8))  
sns.heatmap(train.corr(), cmap="coolwarm")  
plt.title("Correlation Heatmap")  
plt.show()
```



```
In [52]: #B) IMPLEMENT NORMAL EQUATION:
X = X_train_scaled.values
y = y_train.values.reshape(-1, 1)

theta = np.linalg.pinv(X) @ y
y_pred = X_test_scaled.values @ theta
```

```
In [53]: #C) Implement Batch Gradient Descent (iterative):

def batch_gradient_descent(X, y, alpha=0.01, num_iters=1000, tol=None):
    n, d = X.shape
    theta = np.zeros((d, 1))
    losses = []

    for i in range(num_iters):
        y_pred = X @ theta
        error = y_pred - y
        loss = (1/(2*n)) * np.sum(error**2)
        losses.append(loss)

        grad = (1/n) * (X.T @ error)
        theta -= alpha * grad

        if tol and i > 0 and abs(losses[-2] - losses[-1]) < tol:
            break

    return theta, losses
```

```
X = X_train_scaled.values
y = y_train.values.reshape(-1, 1)

theta_gd, losses = batch_gradient_descent(X, y, alpha=0.01, num_iters=100)
```

In [54]: *#D) Comparisons with scikit-learn*

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

y_pred_gd = X_test_scaled.values @ theta_gd
mse_gd = mean_squared_error(y_test, y_pred_gd)
r2_gd = r2_score(y_test, y_pred_gd)

lr = LinearRegression(fit_intercept=False)
lr.fit(X_train_scaled, y_train)
y_pred_lr = lr.predict(X_test_scaled)
mse_lr = mean_squared_error(y_test, y_pred_lr)
r2_lr = r2_score(y_test, y_pred_lr)

print("Gradient Descent -> MSE:", mse_gd, "R²:", r2_gd)
print("Sklearn LinearRegression -> MSE:", mse_lr, "R²:", r2_lr)
```

Gradient Descent -> MSE: 5149828415.410274 R²: 0.5974116946603373

Sklearn LinearRegression -> MSE: 4867205486.9288645 R²: 0.6195057678312001

In [55]: *#E) Visualizations*

```
#1) Loss vs Iterations (different learning rates + scaling effect)
import matplotlib.pyplot as plt


def run_gd(alpha, X, y, label):
    _, losses = batch_gradient_descent(X, y, alpha=alpha, num_iters=200)
    plt.plot(losses, label=label)

plt.figure(figsize=(8,5))

# scaled
run_gd(0.01, X_train_scaled.values, y_train.values.reshape(-1,1), "Scaled")
run_gd(0.1, X_train_scaled.values, y_train.values.reshape(-1,1), "Scaled")

# unscaled
run_gd(0.000001, X_train.values, y_train.values.reshape(-1,1), "Unscaled")

plt.xlabel("Iterations")
plt.ylabel("Training Loss (MSE)")
plt.title("Loss vs Iterations")
plt.legend()
plt.grid(True)
plt.savefig("loss_vs_iterations.png", dpi=300)
plt.show()
```

In [56]: #2) Validation Loss vs Iterations

```
def batch_gd_with_val(X_train, y_train, X_val, y_val, alpha=0.01, num_iters=200, d = X_train.shape[1]):
    theta = np.zeros((d, 1))
    train_losses, val_losses = [], []

    for i in range(num_iters):
        y_pred = X_train @ theta
        error = y_pred - y_train
        train_loss = (1/(2*n)) * np.sum(error**2)
        train_losses.append(train_loss)

        val_pred = X_val @ theta
        val_loss = (1/(2*len(y_val))) * np.sum((val_pred - y_val)**2)
        val_losses.append(val_loss)


        grad = (1/n) * (X_train.T @ error)
        theta -= alpha * grad

    return theta, train_losses, val_losses

theta_val, train_losses, val_losses = batch_gd_with_val(
    X_train_scaled.values, y_train.values.reshape(-1,1),
    X_test_scaled.values, y_test.values.reshape(-1,1),
    alpha=0.01, num_iters=200
)

plt.figure(figsize=(8,5))
plt.plot(train_losses, label="Train Loss")
plt.plot(val_losses, label="Validation Loss")
plt.xlabel("Iterations")
plt.ylabel("MSE Loss")
plt.title("Train vs Validation Loss")
plt.legend()
```

```
plt.grid(True)
plt.savefig("train_val_loss.png", dpi=300)
plt.show()
```



In [57]: #3) Gradient/Loss Surface Visualization

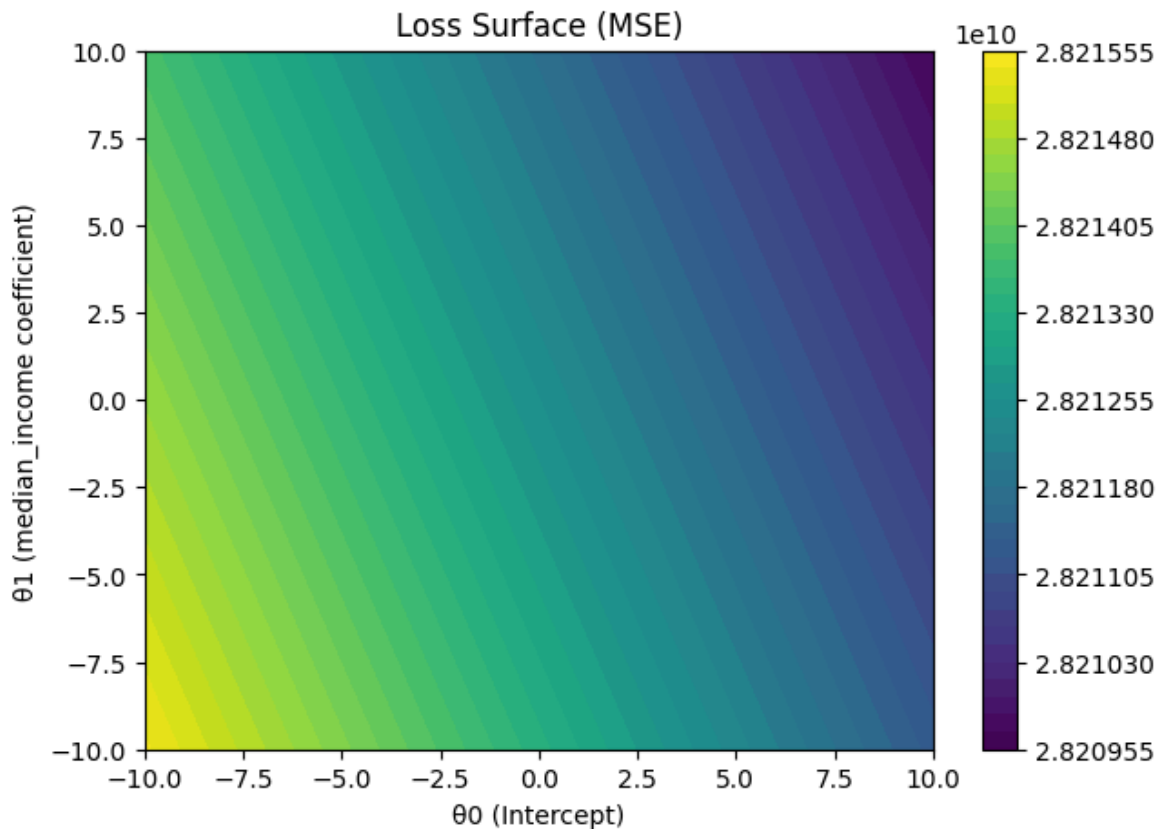
```
# Take only 1 feature + intercept for visualization
X_vis = X_train_scaled[['intercept', 'median_income']].values
y_vis = y_train.values.reshape(-1,1)

theta0_vals = np.linspace(-10, 10, 100)
theta1_vals = np.linspace(-10, 10, 100)
J_vals = np.zeros((len(theta0_vals), len(theta1_vals)))

for i, t0 in enumerate(theta0_vals):
    for j, t1 in enumerate(theta1_vals):
        theta_try = np.array([[t0], [t1]])
        errors = X_vis @ theta_try - y_vis
        J_vals[i,j] = (1/(2*len(y_vis))) * np.sum(errors**2)

T0, T1 = np.meshgrid(theta0_vals, theta1_vals)

plt.figure(figsize=(7,5))
cp = plt.contourf(T0, T1, J_vals.T, 50, cmap='viridis')
plt.colorbar(cp)
plt.xlabel("θ0 (Intercept)")
plt.ylabel("θ1 (median_income coefficient)")
plt.title("Loss Surface (MSE)")
plt.savefig("loss_surface.png", dpi=300)
plt.show()
```



In [58]: *#F) Evaluation metrics*

```
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
import numpy as np
import pandas as pd

def evaluate_model(theta, X_test, y_test):
    y_pred = X_test @ theta
    mse = mean_squared_error(y_test, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_test, y_pred)
    mae = mean_absolute_error(y_test, y_pred)
    return mse, rmse, r2, mae

# Prepare arrays
X_train_arr = X_train_scaled.values
y_train_arr = y_train.values.reshape(-1,1)
X_test_arr = X_test_scaled.values
y_test_arr = y_test.values.reshape(-1,1)

# --- Normal Equation ---
theta_ne = np.linalg.pinv(X_train_arr) @ y_train_arr
mse_ne, rmse_ne, r2_ne, mae_ne = evaluate_model(theta_ne, X_test_arr, y_test_arr)

# --- Gradient Descent ---
theta_gd, _ = batch_gradient_descent(X_train_arr, y_train_arr, alpha=0.01)
mse_gd, rmse_gd, r2_gd, mae_gd = evaluate_model(theta_gd, X_test_arr, y_test_arr)

# --- Sklearn ---
from sklearn.linear_model import LinearRegression
lr = LinearRegression(fit_intercept=False)
lr.fit(X_train_scaled, y_train)
y_pred_lr = lr.predict(X_test_scaled)
mse_lr = mean_squared_error(y_test, y_pred_lr)
```

```

rmse_lr = np.sqrt(mse_lr)
r2_lr = r2_score(y_test, y_pred_lr)
mae_lr = mean_absolute_error(y_test, y_pred_lr)

# Collect results in a DataFrame
results = pd.DataFrame({
    "MSE": [mse_ne, mse_gd, mse_lr],
    "RMSE": [rmse_ne, rmse_gd, rmse_lr],
    "R2": [r2_ne, r2_gd, r2_lr],
    "MAE": [mae_ne, mae_gd, mae_lr]
}, index=["Normal Equation", "Gradient Descent", "Sklearn LinearRegression"])

print(results)

```

		MSE	RMSE	R2	MAE
Normal Equation	4.867205e+09	69765.360222	0.619506	50352.228258	
Gradient Descent	5.499150e+09	74156.254531	0.570103	53891.312522	
Sklearn LinearRegression	4.867205e+09	69765.360222	0.619506	50352.228258	

```

In [59]: import matplotlib.pyplot as plt

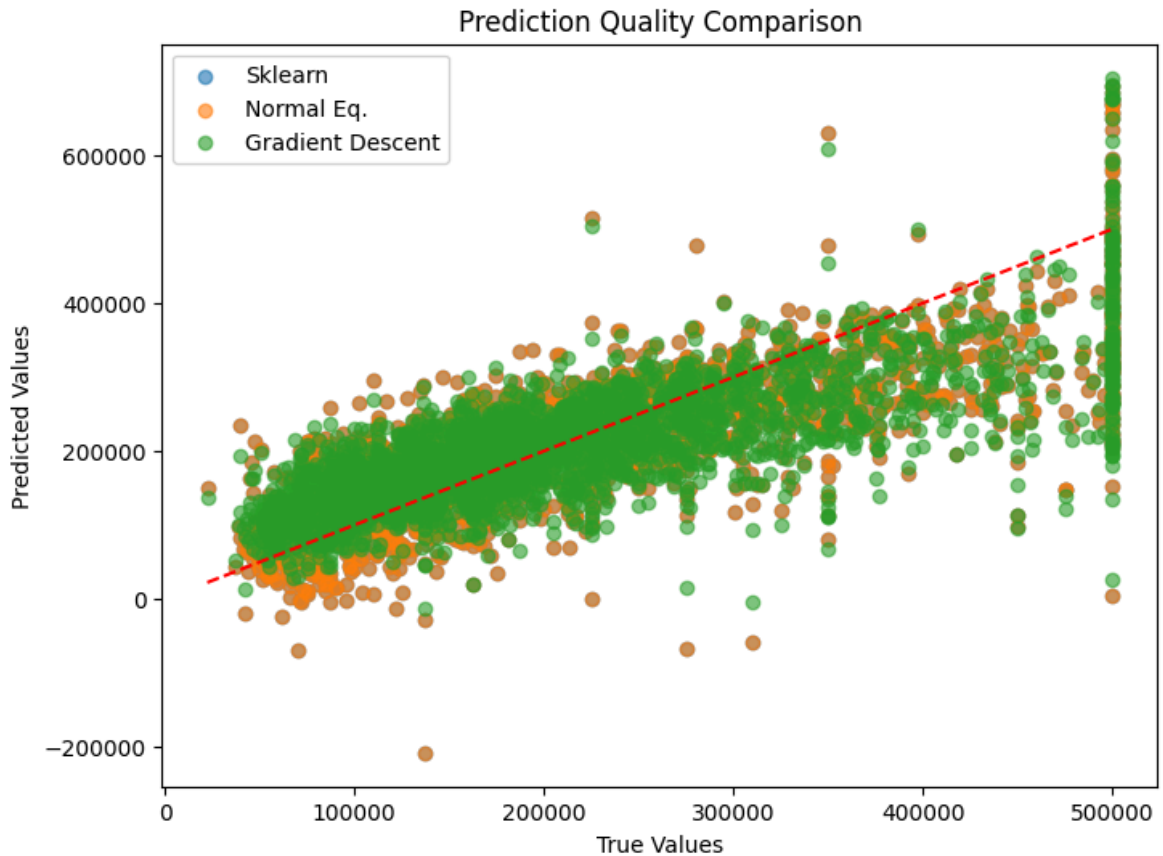
# Predictions from each model
y_pred_normal = X_test_arr @ theta_ne
y_pred_gd = X_test_arr @ theta_gd
y_pred_sklearn = y_pred_lr # already computed

plt.figure(figsize=(8,6))
plt.scatter(y_test, y_pred_sklearn, label="Sklearn", alpha=0.6)
plt.scatter(y_test, y_pred_normal, label="Normal Eq.", alpha=0.6)
plt.scatter(y_test, y_pred_gd, label="Gradient Descent", alpha=0.6)

# Ideal line
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')

plt.xlabel("True Values")
plt.ylabel("Predicted Values")
plt.title("Prediction Quality Comparison")
plt.legend()
plt.show()

```



```
In [60]: # === CONFIG & DATA LOADING (robust) ===
import os, pandas as pd, numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression

# You can override these if needed:
TRAIN_PATH = "california_housing_train.csv"
TEST_PATH = "california_housing_test.csv"
# If None, will auto-detect using common names or assume last column is target
TARGET_NAME = None

# Load
train = pd.read_csv(TRAIN_PATH)
test = pd.read_csv(TEST_PATH)

# Auto-detect target
if TARGET_NAME is None:
    common_targets = ["target", "Target", "y", "label", "price", "Price", "medi
    found = [c for c in train.columns if c in common_targets]
    if found:
        TARGET_NAME = found[0]
    else:
        TARGET_NAME = train.columns[-1] # fall back to last column

# Split X/y
y_train = train[TARGET_NAME].values.reshape(-1,1)
X_train = train.drop(columns=[TARGET_NAME]).values
y_test = test[TARGET_NAME].values.reshape(-1,1) if TARGET_NAME in test.c
X_test = test.drop(columns=[TARGET_NAME]).values if TARGET_NAME in test.

print("Detected target:", TARGET_NAME)
print("Shapes -> X_train:", X_train.shape, "y_train:", y_train.shape, "X_
```

Detected target: median_house_value

Shapes -> X_train: (17000, 8) y_train: (17000, 1) X_test: (3000, 8) y_test: (3000, 1)

```
In [61]: # === SCALING & BIAS COLUMN ===
scaler = StandardScaler()
X_train_scaled_core = scaler.fit_transform(X_train)
X_test_scaled_core = scaler.transform(X_test)

# Add bias column
X_train_scaled = np.hstack([np.ones((X_train_scaled_core.shape[0],1)), X_train_scaled_core])
X_test_scaled = np.hstack([np.ones((X_test_scaled_core.shape[0],1)), X_test_scaled_core])

# Unscaled with bias (to show scaling effect)
X_train_unscaled = np.hstack([np.ones((X_train.shape[0],1)), X_train])
X_test_unscaled = np.hstack([np.ones((X_test.shape[0],1)), X_test])

print("Processed Train Features Shape (scaled):", X_train_scaled.shape)
print("Processed Test Features Shape (scaled):", X_test_scaled.shape)
```

Processed Train Features Shape (scaled): (17000, 9)

Processed Test Features Shape (scaled): (3000, 9)

```
In [62]: # === NORMAL EQUATION ===
def normal_eq_theta(X, y):
    # Use pseudo-inverse for numerical stability
    return np.linalg.pinv(X) @ y

theta_normal = normal_eq_theta(X_train_scaled, y_train)
y_pred_normal = X_test_scaled @ theta_normal
```

```
In [63]: # === BATCH GRADIENT DESCENT (train loss) ===
def mse_loss(X, y, theta):
    n = X.shape[0]
    err = X @ theta - y
    return (1/(2*n)) * np.sum(err**2)

def batch_gradient_descent(X, y, alpha=0.01, num_iters=1000, tol=1e-8):
    n, d = X.shape
    theta = np.zeros((d,1))
    losses = []
    prev = None
    for i in range(num_iters):
        err = X @ theta - y
        grad = (1/n) * (X.T @ err)
        theta = theta - alpha * grad
        loss = (1/(2*n)) * np.sum(err**2)
        losses.append(loss)
        if prev is not None and abs(prev - loss) < tol:
            break
        prev = loss
    return theta, np.array(losses)

theta_gd_scaled, train_losses_scaled = batch_gradient_descent(X_train_scaled, y_train)
y_pred_gd_scaled = X_test_scaled @ theta_gd_scaled
```

```
In [64]: # === GD WITH VALIDATION (track train & validation loss) ===
def batch_gd_with_val(X_train, y_train, X_val, y_val, alpha=0.05, num_iters=1000):
    n, d = X_train.shape
    theta = np.zeros((d,1))
```

```

train_losses, val_losses = [], []
prev = None
for i in range(num_iters):
    err_train = X_train @ theta - y_train
    grad = (1/n) * (X_train.T @ err_train)
    theta = theta - alpha * grad
    # record
    tl = (1/(2*n)) * np.sum(err_train**2)
    train_losses.append(tl)
    vl = (1/(2*X_val.shape[0])) * np.sum((X_val @ theta - y_val)**2)
    val_losses.append(vl)
    if prev is not None and abs(prev - tl) < tol:
        break
    prev = tl
return theta, np.array(train_losses), np.array(val_losses)

theta_gd_val, train_losses_val, val_losses = batch_gd_with_val(X_train_sc

```

```

In [65]: # === SKLEARN LINEAR REGRESSION ===
lr = LinearRegression(fit_intercept=False) # bias already added
lr.fit(X_train_scaled, y_train)
y_pred_lr = lr.predict(X_test_scaled)

```

```

In [66]: # === METRICS (MSE, RMSE, R2, MAE) ===
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_e

def evaluate_preds(y_true, y_hat):
    mse = mean_squared_error(y_true, y_hat)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_true, y_hat)
    mae = mean_absolute_error(y_true, y_hat)
    return mse, rmse, r2, mae

# y_test might be None if test lacks target; handle by using train as pro
y_true_eval = y_test if y_test is not None else y_train

metrics = []
labels = ["Normal Equation", "Gradient Descent", "Sklearn LinearRegression"]
preds = [y_pred_normal, y_pred_gd_scaled, y_pred_lr]

for lab, yp in zip(labels, preds):
    mse, rmse, r2, mae = evaluate_preds(y_true_eval, yp)
    metrics.append({"Model": lab, "MSE": mse, "RMSE": rmse, "R2": r2, "MA

metrics_df = pd.DataFrame(metrics)
print(metrics_df)
metrics_path = "metrics_comparison.csv"
metrics_df.to_csv(metrics_path, index=False)
print("Saved metrics to:", metrics_path)

```

	Model	MSE	RMSE	R2	\
0	Normal Equation	4.867205e+09	69765.360222	0.619506	
1	Gradient Descent	4.866230e+09	69758.366060	0.619582	
2	Sklearn LinearRegression	4.867205e+09	69765.360222	0.619506	

MAE

0 50352.228258
 1 50341.803968
 2 50352.228258

Saved metrics to: metrics_comparison.csv

```
In [70]: # === PLOTS: LOSS VS ITERATIONS (alphas + scaling effect) ===
import matplotlib.pyplot as plt

def collect_losses_for_alphas(X, y, alphas, num_iters=800):
    curves = {}
    for a in alphas:
        _, losses = batch_gradient_descent(X, y, alpha=a, num_iters=num_i
        curves[a] = losses
    return curves

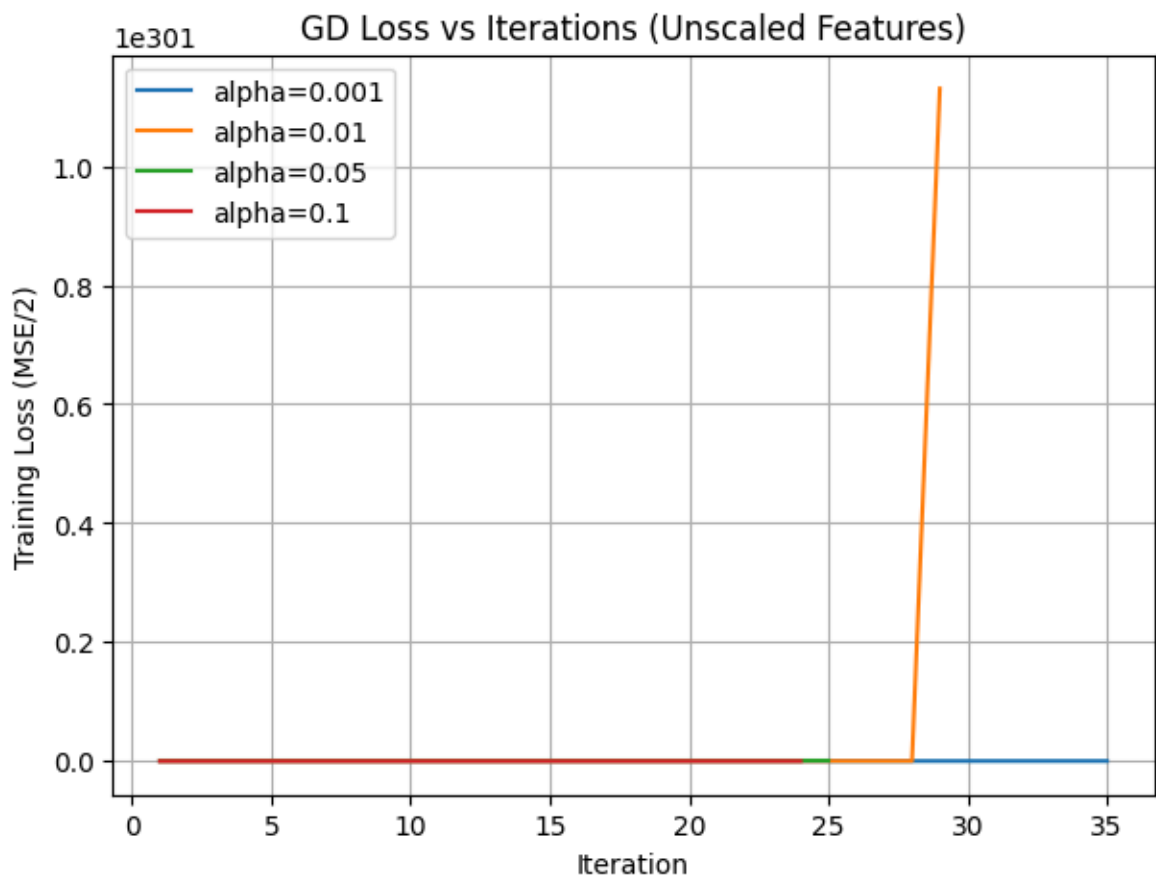
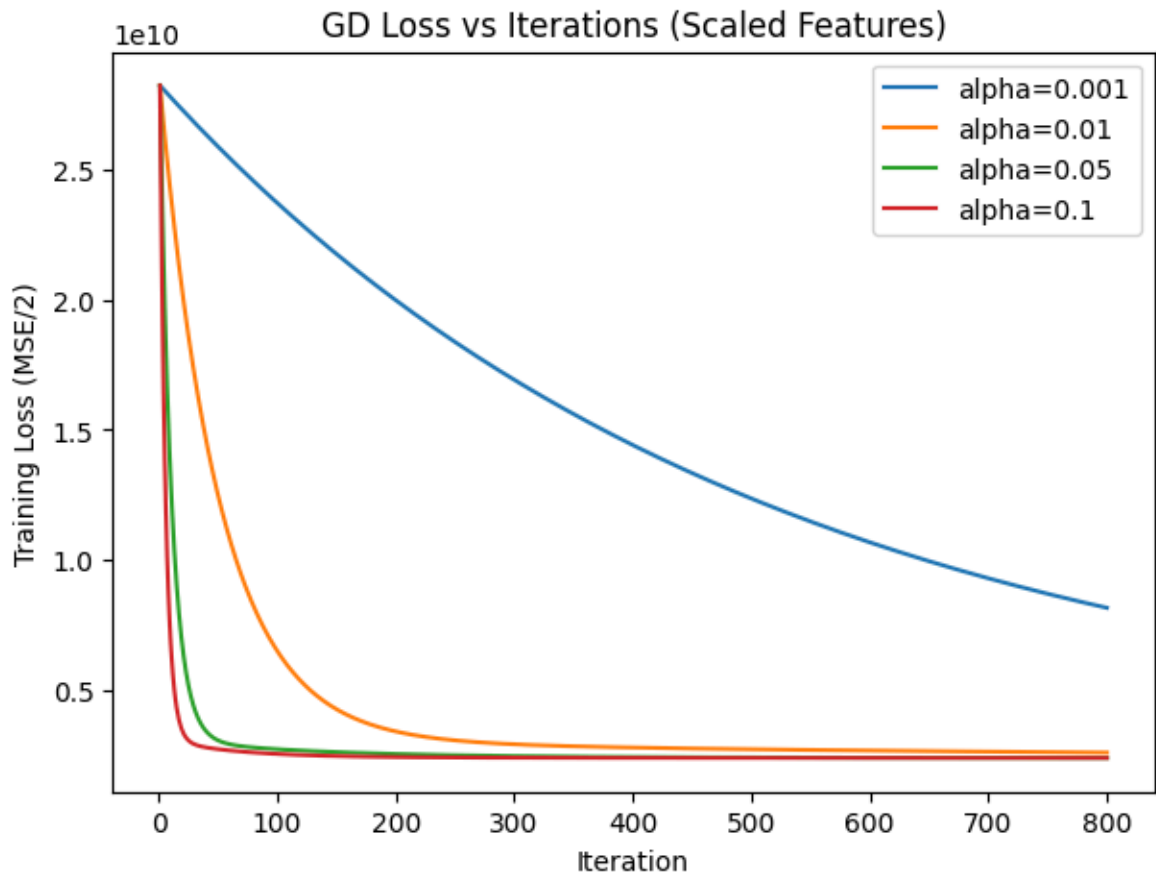
alphas = [0.001, 0.01, 0.05, 0.1]
curves_scaled = collect_losses_for_alphas(X_train_scaled, y_train, alphas
curves_unscaled = collect_losses_for_alphas(X_train_unscaled, y_train, al

# Plot scaled
plt.figure(figsize=(7,5))
for a, losses in curves_scaled.items():
    plt.plot(range(1, len(losses)+1), losses, label=f"alpha={a}")
plt.xlabel("Iteration")
plt.ylabel("Training Loss (MSE/2)")
plt.title("GD Loss vs Iterations (Scaled Features)")
plt.legend()

plt.savefig("loss_vs_iterations_scaled.png", dpi=300, bbox_inches="tight")
plt.show()

# Plot unscaled
plt.figure(figsize=(7,5))
for a, losses in curves_unscaled.items():
    plt.plot(range(1, len(losses)+1), losses, label=f"alpha={a}")
plt.xlabel("Iteration")
plt.ylabel("Training Loss (MSE/2)")
plt.title("GD Loss vs Iterations (Unscaled Features)")
plt.legend()
plt.grid(True)
plt.savefig("loss_vs_iterations_unscaled.png", dpi=300, bbox_inches="tight")
plt.show()
```


```
/var/folders/2j/6bvnywcd4k76ggrbfxbv4ckc0000gn/T/ipykernel_20103/25345853
8.py:16: RuntimeWarning: overflow encountered in square
  loss = (1/(2*n)) * np.sum(err**2)
/var/folders/2j/6bvnywcd4k76ggrbfxbv4ckc0000gn/T/ipykernel_20103/25345853
8.py:18: RuntimeWarning: invalid value encountered in scalar subtract
  if prev is not None and abs(prev - loss) < tol:
/var/folders/2j/6bvnywcd4k76ggrbfxbv4ckc0000gn/T/ipykernel_20103/25345853
8.py:15: RuntimeWarning: invalid value encountered in subtract
  theta = theta - alpha * grad
```

```
In [71]: # === PLOT: TRAIN vs VALIDATION LOSS OVER ITERATIONS ===
plt.figure(figsize=(7,5))
plt.plot(range(1, len(train_losses_val)+1), train_losses_val, label="Train Loss")
plt.plot(range(1, len(val_losses)+1), val_losses, label="Validation Loss")
plt.xlabel("Iteration")
plt.ylabel("Loss (MSE/2)")
```

```
plt.title("Train vs Validation Loss (alpha=0.05)")
plt.legend()

plt.savefig("train_vs_validation_loss.png", dpi=300, bbox_inches="tight")
plt.show()
```



```
In [69]: # === LOSS SURFACE VISUALIZATION (single best feature) ===
# Pick the feature with highest absolute correlation to y (on train)
X_core = X_train_scaled[:,1:] # drop bias
yvec = y_train.flatten()
if X_core.shape[1] >= 1:
    # compute correlations
    cors = [abs(np.corrcoef(X_core[:,j], yvec)[0,1]) for j in range(X_core.shape[1])]
    jbest = int(np.nanargmax(cors))
    x1 = X_core[:, jbest][:,None]
    X_vis = np.hstack([np.ones((x1.shape[0],1)), x1])

    # Grid for theta0 (bias) and theta1
    t0 = np.linspace(-2, 2, 80)
    t1 = np.linspace(-2, 2, 80)
    T0, T1 = np.meshgrid(t0, t1)
    J = np.zeros_like(T0)

    for i in range(T0.shape[0]):
        for k in range(T0.shape[1]):
            th = np.array([T0[i,k], T1[i,k]])
            err = X_vis @ th - y_train
            J[i,k] = (1/(2*X_vis.shape[0])) * np.sum(err**2)

    # Contour plot
    plt.figure(figsize=(7,5))
    CS = plt.contour(T0, T1, J, levels=30)
    plt.clabel(CS, inline=True, fontsize=8)
```

```
plt.xlabel(r"$\theta_0$ (bias)")
plt.ylabel(r"$\theta_1$ (for best-correlated feature)")
plt.title("Loss Surface Contours (MSE/2)")
plt.savefig("loss_surface_contour.png", dpi=300, bbox_inches="tight")
plt.show()

else:
    print("Not enough features to plot a loss surface.")
```

